

Song Search using CLAP

SUSE Hack Week 25



Goals

Search music using natural language descriptions

- Find songs by describing what you hear: "piano melody", "female vocalist"
- Compare songs to find similar tracks
- Experiment AI with Copilot Code Assistant

CLAP Overview

CLAP (Contrastive Language-Audio Pretraining) - Bridging audio and language

Core Concept:

CLAP can represent both songs and text descriptions as **512-dimensional vectors** in the same mathematical space. This enables searching for music using natural language queries.

How it works:

- A song becomes a 512-number vector capturing its audio characteristics
- A text query like "piano music" also becomes a 512-number vector
- Vectors that point in similar directions represent similar content
- We measure similarity to find songs matching the description

Model: music_audioset_epoch_15_esc_90.14.pt - Specialized for music (71% genre accuracy, trained on ~4M samples)



Implementation

Two Modes:

- **Single Analysis:** `single_analysis.py` - One song vs one query
- **Batch Analysis:** `multiple_analysis.py` - All songs vs all queries

Architecture:

- `clap_analysis.py` - Core library (shared)
- Automatic model download & caching
- Parallel processing across CPU cores
- CSV output with detailed metrics

Cosine Similarity Scores:

- Range: -1 (opposite) to $+1$ (identical)
- Thresholds: $> 0.3 = \text{HIGH}$, $> 0.15 = \text{MODERATE}$, $\leq 0.15 = \text{LOW}$

Output Example - Classical

Song: Aaron Dunn - Minuet (Classical Piano)

| HIGH (> 0.3) | MODERATE (> 0.15) | LOW (≤ 0.15) |
|---------------------------------------------|-----------------------|--------------------------------|
| 0.41 piano | 0.29 chill | 0.15 70s |
| 0.40 classical instrumental song with piano | 0.28 chillout | 0.14 choir |
| 0.38 classical music | 0.27 sad | 0.14 beautiful |
| 0.37 instrumental | 0.26 jazz | 0.14 danceable |
| 0.33 instrumental music | 0.25 acoustic | 0.12 pop music wit bass guitar |
| | 0.23 acoustic guitar | 0.12 funk |
| | 0.22 pop | 0.12 bass guitar |
| | 0.21 folk | 0.11 soul |
| | 0.21 happy | 0.11 rock |

Observation: Piano and classical-related queries score highest. Chill/mellow descriptors moderate. Rock and modern era queries score low.

Tip: Detailed, specific queries combining multiple descriptors (genre + instrument + style) yield excellent results!

Output Example - Rock

Song: SUSE Band - Paint it Green (Rock)

| HIGH (> 0.3) | MODERATE (> 0.15) | LOW (≤ 0.15) |
|-----------------------|-------------------|----------------------|
| 0.42 indie rock | 0.28 hard rock | 0.14 heavy metal |
| 0.41 alternative rock | 0.28 folk | 0.14 dance |
| 0.40 classic rock | 0.26 punk | 0.13 soul |
| 0.37 Progressive rock | 0.25 80s | 0.13 instrumental |
| 0.37 indie pop | 0.25 funk | 0.13 acoustic guitar |
| 0.36 rock | 0.2590s | 0.13 singing |
| 0.33 female vocalist | 0.2460s | 0.12 bass guitar |
| 0.32 rock music | 0.24 metal | 0.12 guitar |
| 0.31 blues | 0.23 Mellow | 0.06 female voice |

Observation: Rock subgenres (indie, alternative, classic) score highest. Generic instrument queries score low.

Query Wording Matters: Notice "female vocalist" (0.33, HIGH) vs "female voice" (0.06, LOW) - 5x difference! Use music industry vocabulary for better results.

Finding Similar Songs

Bonus Feature: Beyond text search, we can find which songs sound similar to each other.

Example: Finding songs similar to SUSE Band - Paint it Green (Rock)

| Song | Similarity |
|----------------------------------------|------------|
| HoliznaCC0 - Dreams Of Lilith (Rock) | 0.72 |
| Jenna Jay - Someone Real (Pop) | 0.56 |
| Zane Little - Always and Forever (Pop) | 0.47 |
| SUSE Band - SUSE Yes Please (Pop) | 0.43 |
| SUSE Band - Can't Stop The SUSE (Pop) | 0.39 |
| Aaron Dunn - Minuet (Classical Piano) | 0.18 |

Use Cases:

- Find songs with similar instrumentation or mood
- Create automatic playlists based on audio similarity
- Discover patterns in your music collection

Containerized Web App

SUSE Hack week 25 - Song Search

Music Search & Discovery

Songs Analyzed: 52 Music Library: 2h 53m Queries: 1

Analysis Text Search Similar Songs

Text Search

Search songs using natural language descriptions.

Search

Top 20 Results

| # | Artist | Title | Score | Play |
|---|-------------------------------|--------------------------------------|--------|------|
| 1 | Madoka Ogitani | Ibuki | 0.4502 | Play |
| 2 | 06 HoliznaCC0 | Spring On The Horizon | 0.3894 | Play |
| 3 | Aaron Dunn | Minuet - Notebook for Anna Magdalena | 0.3501 | Play |
| 4 | Daniele Pasini/Raffaele Pilia | "Memoria d'azzurro" | 0.3376 | Play |

AI Support in Development - What Went Well

Project developed with GitHub Copilot (Claude Sonnet 4.5)

1. Research & Model Selection

- Compared CLAP models → selected music_audioset (71% GTZAN accuracy)
- Read CLAP paper, explained architecture accessibly

2. Code Analysis & Technical Discovery

- Examined checkpoint: model lacks fusion, max 10s segments
- Designed overlapping segment solution

3. Data Analysis & Documentation

- Found "vocalist" vs "voice" pattern (5x difference in results)
- Generated all technical docs with examples

1. Unrequested Code Changes

- Single-file CLI disappeared, score icons changed without asking
- **Problem:** Full file rewrites hide unintended changes
- **Lesson:** Always review diffs carefully

2. First Solution \neq Best Solution

- **Performance:** Song analyzed per query \rightarrow challenged \rightarrow **15x faster**
- **Threading:** Single-threaded \rightarrow asked "faster?" \rightarrow multiprocessing

Key Lesson: Working \neq optimal. Challenge AI with domain knowledge.

Best Practices for AI-Assisted Development

Do's:

- Ask AI to research and compare options (models, approaches)
- Request code analysis to understand existing implementations
- Share actual data/errors for pattern recognition
- Challenge solutions: "Can this be faster?" "Is there a better way?"
- Request documentation with examples from your project

Don'ts:

- Don't accept first working solution as optimal
- Don't let AI rewrite entire files without showing diffs
- Don't assume AI made only requested changes - review carefully
- Don't trust performance without testing alternatives

Bottom line: AI is a powerful assistant for research, analysis, and documentation. For code: it provides a working starting point, but optimization requires your expertise and critical thinking.



References

- **CLAP:** <https://github.com/LAION-AI/CLAP> - The main model being researched
- **Hugging Face:**
https://huggingface.co/docs/transformers/model_doc/clap
 - Pre-trained models for CLAP
- **Free Music Archive:** <https://freemusicarchive.org/> - Creative Commons songs for testing
- **SUSE Hack Week 25 Project:**
<https://hackweek.opensuse.org/25/projects/clap-machine-learning-to-search-song-starting-from-text>
 - Project page
- **Project Repository:**
<https://github.com/gcolangiulisuse/hw25-song-search> - GitHub repo of the project